


COMP3702/COMP7702 Artificial Intelligence

Semester 2, 2020

Tutorial 11

Example domain

Consider the gridworld below, which we saw in Tutorial 9:



States in this environment are the positions on the tiles. The world is bounded by a boundary wall, and there is one obstacle, at $[1,1]$ (using python or zero indexing starting from the bottom left corner). In addition, there are two terminal states, indicated by the coloured squares.

Actions and Transitions: In this world, an agent can generally choose to move in four directions — *up*, *down*, *left* and *right*. However, the agent moves successfully with only $p = 0.8$, and moves perpendicular to its chosen direction with $p = 0.1$ in each perpendicular direction. If it hits a wall or obstacle, the agent stays where it is (i.e. no collision cost). In addition, once the agent arrives on a coloured square with a value, it has only one special action available to it; that is, to *exit* the environment.

Rewards: The values stated on the coloured squares are the reward for *exiting* the square and the environment, so the reward is not repeatedly earned; that is, $R([3,2], \text{exit}) = a$, and $R([3,1], \text{exit}) = b$. All other states have 0 reward.

Discount factor: $\gamma = 0.9$.

Exercise 11.1.

- List the dimensions of complexity for a reinforcement learning environment: Modularity, Planning horizon, Representation, Computational limits, Learning, Sensing uncertainty, Effect uncertainty, Preference, Number of agents, Interaction.
- In RL, what is the connection between the environment's state transition function and the exploration policy used by the agent?

Exercise 11.2.

For the grid world above, develop a simulator for the state-action-reward-state loop.

- In particular, you should develop a `Grid.apply_move()` function that takes the current state and actions as inputs and updates the state while returning the reward for the state-action-state transition.
- The provided code includes `Grid.player_x` and `Grid.player_y` variables that represent the state of the player on the (x, y) grid.
- You should make use of the provided code, especially the `stoch_action()` and `attempt_move()` methods.
- Include a way to restart the agent in a random map location, avoiding obstacles, after it exits the environment.

(Note: you may wish to look at how the Laser Tank environment's `apply_move()` method works, as provided in the Assignment 4 support code).

Exercise 11.3. Using your simulator, implement Q-learning for this gridworld problem.

- a) First, write a function to choose an action given the Q-value estimates. Incorporate your agent's exploration strategy in this function.
- b) Then write a `next_iteration()` function, which will probe the simulator to receive a reward and a new state, and use this to update your agent's Q-values.
 - Be sure to include an exploration strategy
 - See the provided code.